Data-Driven MLOps Pipeline for Anomaly Detection

Enterprise-Grade Automated Machine Learning Deployment System

Domain	Machine Learning Operations (MLOps) · DevSecOps · Cloud-Native Engineering
Technology	Python · FastAPI · IsolationForest · Docker · Kubernetes · Jenkins · Prometheus · Grafana
Version	1.0.0
Date	May 2026

[Python 3.10+] [FastAPI 0.111] [Scikit-Learn 1.4] [Docker] [Minikube K8s] [Jenkins LTS] [Prometheus] [Grafana]

> ***Project Title:*** *Data-Driven MLOps Pipeline for Real-Time Fraud Anomaly Detection*
>
> ***Technology Domain:*** *Machine Learning Operations (MLOps) · DevSecOps · Cloud-Native Engineering*
>
> ***Version:*** *1.0.0 · May 2026*

Table of Contents

---

1. Introduction

1.1 Background

The financial services industry loses an estimated **$32 billion annually** to payment card fraud globally (Nilson Report, 2023). Traditional rule-based fraud detection systems, while interpretable, suffer from a fundamental limitation: they require human analysts to manually define thresholds and decision rules that fraudsters can study, reverse-engineer, and evade over time. Machine learning offers an alternative — pattern recognition directly from transaction data without explicit rule definition.

However, the gap between a Data Scientist producing a trained model on a laptop and that model serving production traffic reliably, securely, and at scale represents one of the most significant unsolved challenges in applied AI. This gap is the domain of **Machine Learning Operations (MLOps)**.

1.2 Project Overview

This project implements a complete, end-to-end **MLOps pipeline for real-time transaction fraud detection**. It takes an unsupervised anomaly detection model (Scikit-Learn `IsolationForest`) and operationalizes it into a production-grade, containerized microservice that is:

- **Automatically built and tested** on every code commit via a Jenkins CI/CD pipeline
- **Security-scanned** for vulnerabilities before any artifact is deployed
- **Containerized** into a minimal, hardened Docker image following the principle of least privilege
- **Orchestrated** on a local Kubernetes cluster (Minikube) with high-availability, self-healing, and auto-scaling
- **Monitored** in real-time through Prometheus metric collection and Grafana dashboards

The project intentionally bridges two disciplines — **Data Science** (model development, feature engineering, serialization) and **Advanced DevOps** (CI/CD, containerization, Kubernetes, observability) — to demonstrate how modern MLOps unifies them into a single, automated lifecycle.

1.3 Objectives

#	Objective
1	Train and serialize an `IsolationForest` anomaly detection model on synthetic transaction data
2	Expose the model as a RESTful microservice via FastAPI
3	Implement a DevSecOps-compliant Jenkins CI/CD pipeline with 7 automated stages
4	Containerize the application using a production-optimized, multi-stage Docker build
5	Deploy and orchestrate the service on Minikube (local Kubernetes) with HA configuration
6	Establish full observability with Prometheus telemetry and Grafana visualization

1.4 Scope

The scope of this project covers the full MLOps lifecycle from model training through production monitoring within a local, Windows-based development environment using Docker Desktop and Minikube. The system is architected using industry-standard open-source tooling and follows patterns directly applicable to cloud-native production deployments (AWS EKS, GCP GKE, Azure AKS).

---

2. Profile of the Problem

2.1 Rationale

Financial fraud is not a static problem. The profile of fraudulent transactions evolves continuously as fraud actors adapt to detection methods. Two features are particularly diagnostic across academic literature and industry datasets:

1. **Transaction Amount** — Fraudulent transactions tend to cluster in the high-value tail of the distribution. A stolen card is typically used for maximum-value purchases before the card is blocked.
2. **Distance from Home** — Genuine cardholders have predictable geographic patterns. A transaction 1,200 km from a cardholder's registered address, especially combined with a high amount, is a strong anomaly signal.

The core statistical challenge is that **fraud labels are almost never available in real-time**. Labelling requires post-hoc investigation — a process that takes days to weeks. This makes **unsupervised learning** the natural approach: the model learns the shape of normal behaviour and flags deviations, without requiring any labelled fraud examples.

2.2 Problem Statement

> ***"How can a financial institution continuously deploy, monitor, and self-heal a machine learning fraud detection model in production — with full automation, security scanning, and zero-downtime deployments — without requiring manual intervention for routine operations?"***

The specific sub-problems addressed are:

Sub-Problem	Challenge
Model Serving	A `.pkl` file on disk must become a low-latency, scalable HTTP API

Sub-Problem	Challenge
Automation	Each code change must trigger a full build-test-deploy cycle without human steps
Security	Source code and container images must be scanned before reaching production
Resilience	The service must self-heal from pod failures and auto-scale under load
Observability	Operations teams must be able to see prediction rates, latency, and anomaly scores in real time

2.3 Scope of Study

In Scope:

- Unsupervised anomaly detection using Isolation Forest on tabular transaction features
- REST API microservice development with FastAPI
- CI/CD pipeline implementation with Jenkins (7 stages)
- Docker containerization (multi-stage, security-hardened)
- Kubernetes deployment on Minikube (Deployment, Service, HPA)
- Prometheus metric collection and Grafana dashboard integration
- Comprehensive automated testing (unit, integration, boundary)

Out of Scope:

- Supervised fraud classification (requires labelled datasets)
- Cloud provider deployments (AWS/GCP/Azure)
- Real payment card dataset ingestion (PCI-DSS compliance)
- User authentication and API key management
- Alertmanager configuration for PagerDuty/Slack notifications

3. Existing System

3.1 Introduction to Existing Approaches

Before modern MLOps pipelines, financial institutions relied on two primary approaches to fraud detection:

a) Rule-Based Expert Systems

Analysts manually encode fraud heuristics as boolean rules: "flag if amount > $2,000 AND country != registered country". Vendors such as FICO (Falcon Fraud Manager) and IBM Financial Crimes Insight offer configurable rule engines. These systems are highly interpretable but brittle — rules must be manually updated as fraud patterns evolve and they generate high false-positive rates.

b) Batch Machine Learning Pipelines

Periodic (nightly or weekly) retraining pipelines using tools like Apache Spark MLlib or SAS Analytics. Models are trained offline, evaluated by a data scientist, and manually promoted to production by a deployment team. This process introduces days of latency between model improvement and production deployment, and the handoff is entirely manual.

3.2 Existing Software in Common Use

Tool/System	Category	Limitation
FICO Falcon	Rule-based fraud detection	Requires manual rule updates; proprietary

Tool/System	Category	Limitation
IBM OpenPages	Risk management platform	No real-time ML inference; governance-only
Apache Spark MLlib	Batch ML pipeline	Not designed for real-time serving
AWS SageMaker	Cloud MLOps	Vendor lock-in; costly; not local
MLflow (standalone)	Experiment tracking	No built-in serving or CI/CD integration
Kubeflow	Kubernetes ML pipelines	Very high operational complexity for initial setup

3.3 Data Flow Diagram — Present (Manual) System

3.4 What Is New in the System Being Developed

The proposed system eliminates every manual step in the above pipeline and adds capabilities not present in any existing tool:

Feature	Existing System	New System
Model-to-API promotion	Manual (days)	Automated via Jenkins (minutes)
Security scanning	None	Bandit SAST with HIGH-severity gate

Feature	Existing System	New System
Testing gate	Manual / ad hoc	30 automated pytest cases in CI
Container build	Manual `docker build`	Jenkins stage with build number tagging
Deployment	Manual `kubectl apply`	Automated with `rollout status` wait
Rollback	Manual	Automatic `kubectl rollout undo` on failure
Observability	Log inspection	Prometheus counters + Grafana dashboards
Self-healing	Manual pod restart	Kubernetes liveness/readiness probes
Auto-scaling	Manual replica adjustment	HPA (2→5 replicas on CPU/memory)
Zero-downtime deploy	Not guaranteed	Rolling update (maxUnavailable: 0)

4. Problem Analysis

4.1 Product Definition

Product Name: MLOps Fraud Detection Pipeline

Product Type: Enterprise-grade MLOps microservice with full CI/CD automation

Primary Users: DevOps/MLOps Engineers, Data Scientists, Security Engineers

The product is a system of interconnected components, not a single application. Its value lies in the **pipeline** — the automated chain of processes that takes a model from code commit to monitored production deployment.

Core Components:

1. `app/model.py` — ML training and serialization engine
2. `app/main.py` — FastAPI inference microservice
3. `Dockerfile` — Multi-stage, security-hardened container build
4. `Jenkinsfile` — 7-stage declarative CI/CD automation pipeline
5. `deployment.yaml` — Kubernetes orchestration manifests (Deployment + Service + HPA)
6. `config/prometheus.yml` — Observability scrape configuration

4.2 Feasibility Analysis

4.2.1 Technical Feasibility

All tools used are open-source, actively maintained, and widely adopted in enterprise environments:

Component	Tool	Maturity
ML Framework	Scikit-Learn 1.4.2	Very High — industry standard since 2007
API Framework	FastAPI 0.111	High — 70k+ GitHub stars, production-proven
Container Engine	Docker 4.x	Very High — de facto container standard

Component	Tool	Maturity
Orchestration	Kubernetes / Minikube	Very High — dominant orchestration platform
CI/CD	Jenkins LTS	Very High — 20+ years enterprise CI leader
Monitoring	Prometheus + Grafana	Very High — CNCF graduated projects

The system runs entirely on a single Windows machine with Docker Desktop and Minikube, requiring:

- Minimum 8 GB RAM (4 GB for Minikube, 2 GB for Jenkins, 2 GB for host)
- Minimum 4 CPU cores
- 20 GB free disk space

4.2.2 Economic Feasibility

Total licensing cost: **$0** — the entire stack is open-source.

Operational costs are limited to electricity and hardware depreciation for the development machine. In a production cloud environment, the equivalent stack would cost approximately $200–500/month on a major cloud provider, still representing significant savings over proprietary MLOps platforms ($5,000–50,000/month).

4.2.3 Operational Feasibility

The system is designed for minimal operational overhead:

- Jenkins automates all deployment tasks triggered by Git push
- Kubernetes self-heals failed pods automatically
- HPA manages scaling without human intervention
- Grafana dashboards provide at-a-glance health status

The primary operational skill required is basic Kubernetes and Docker knowledge, which is standard for any DevOps team.

4.3 Project Plan

```
Week 1: Foundation & Model Development
  ■■■ Environment setup (Python venv, Docker Desktop, Minikube)
  ■■■ IsolationForest model design and hyperparameter selection
  ■■■ app/model.py — training and serialization logic
  ■■■ Initial unit tests for model loading

Week 2: API Layer & Containerization
  ■■■ app/main.py — FastAPI endpoints (/predict, /metrics, /)
  ■■■ Pydantic request/response schema design
  ■■■ Prometheus metrics integration
  ■■■ Multi-stage Dockerfile construction
  ■■■ Comprehensive test_main.py test suite

Week 3: CI/CD Pipeline & Kubernetes
  ■■■ Local Docker registry setup
  ■■■ Jenkins container deployment and configuration
  ■■■ Jenkinsfile — 7-stage declarative pipeline
  ■■■ deployment.yaml — Deployment, Service, HPA manifests
  ■■■ Pipeline integration testing

Week 4: Observability & Documentation
  ■■■ Prometheus configuration and local deployment
  ■■■ Grafana setup and dashboard creation
  ■■■ README.md operational runbook
  ■■■ Project report compilation
```

5. Software Requirement Analysis

5.1 Introduction

This section formally specifies the functional and non-functional requirements of the MLOps Fraud Detection Pipeline. Requirements are expressed using standard notation: **FR** (Functional Requirement) and **NFR** (Non-Functional Requirement), with priority levels **[HIGH]**, **[MEDIUM]**, **[LOW]**.

5.2 General Description

5.2.1 Product Perspective

The system is a standalone MLOps pipeline that operates independently of any external cloud services. It communicates with a local Kubernetes cluster via `kubectl`, stores Docker images in a local registry, and exposes its API on a fixed NodePort.

5.2.2 Product Functions

- Train and serialize an Isolation Forest anomaly detection model
- Serve real-time predictions via REST API
- Automate the full build-test-deploy lifecycle on code commit
- Monitor production performance via Prometheus/Grafana

5.2.3 User Classes

User Class	Description	Primary Interface
Data Scientist	Modifies model training logic	`app/model.py`, CLI
API Consumer	Sends transaction data for scoring	`POST /predict`
DevOps Engineer	Manages pipeline, cluster, monitoring	Jenkins, kubectl, Grafana
Security Auditor	Reviews SAST reports and compliance	Bandit JSON report, Jenkins artifacts

5.2.4 Operating Environment

- **Host OS:** Windows 10/11
- **Container Engine:** Docker Desktop 4.x with WSL2 backend
- **Kubernetes:** Minikube v1.35+ with Docker driver
- **Python Runtime:** CPython 3.10+
- **CI Server:** Jenkins LTS 2.440+ in Docker

5.3 Specific Requirements

5.3.1 Functional Requirements

FR-01 [HIGH] — Model Training

The system SHALL train an `IsolationForest` model with `n_estimators=200`, `contamination=0.04`, and a fixed `random_state=42` for reproducibility. Training data SHALL consist of 10,000 normal and 400 anomalous synthetic transactions.

FR-02 [HIGH] — Artifact Serialization

The system SHALL serialize the trained model (`model.pkl`) and feature scaler (`scaler.pkl`) to disk using `joblib` with compression level 3.

FR-03 [HIGH] — Health Endpoint

`GET /` SHALL return HTTP 200 with `{"status": "healthy"}` when the model is loaded. It SHALL return HTTP 503 when the model is not loaded, to signal Kubernetes probes.

FR-04 [HIGH] — Prediction Endpoint

`POST /predict` SHALL accept a JSON body with `amount` (float, 0 < amount ≤ 1,000,000) and `distance_from_home` (float, 0 ≤ distance ≤ 20,000) and return `is_anomaly` (bool), `anomaly_score` (float), `label` (str), `model_version` (str), and `processing_time_ms` (float).

FR-05 [HIGH] — Input Validation

The system SHALL reject requests with: negative amounts, zero amounts, amounts exceeding $1,000,000, negative distances, distances exceeding 20,000 km, NaN values, Inf values, missing fields, and non-numeric values. Rejection SHALL return HTTP 422.

FR-06 [HIGH] — Prometheus Metrics Endpoint

`GET /metrics` SHALL return Prometheus text exposition format metrics including: total request count (by method/endpoint/status), prediction count (by result), prediction latency histogram, last anomaly score, and model loaded status.

FR-07 [HIGH] — CI Pipeline Checkout

The Jenkins pipeline SHALL check out the latest source code from the configured SCM on every build trigger.

FR-08 [HIGH] — CI Pipeline Security Scan

The Jenkins pipeline SHALL run Bandit SAST on the `app/` directory. The pipeline SHALL fail if any HIGH severity finding is detected. Results SHALL be archived as a build artifact.

FR-09 [HIGH] — CI Pipeline Unit Tests

The Jenkins pipeline SHALL execute `pytest tests/` and publish JUnit XML results. The pipeline SHALL NOT proceed to Docker build if any test fails.

FR-10 [HIGH] — CI Pipeline Docker Build

The Jenkins pipeline SHALL build a Docker image tagged with the `BUILD_NUMBER` and `latest`, using the `--platform linux/amd64` flag.

FR-11 [HIGH] — CI Pipeline Registry Push

The Jenkins pipeline SHALL push both the versioned and `:latest` tags to the local Docker registry at `localhost:5000`.

FR-12 [HIGH] — CI Pipeline Kubernetes Deploy

The Jenkins pipeline SHALL apply Kubernetes manifests, update the container image to the build-number-tagged version, and wait for `kubectl rollout status` to confirm successful deployment.

FR-13 [HIGH] — CI Pipeline Automatic Rollback

If the Kubernetes deployment stage fails, the Jenkins pipeline SHALL automatically execute `kubectl rollout undo` before marking the build as failed.

FR-14 [MEDIUM] — Post-Deploy Smoke Test

The Jenkins pipeline SHALL send a test `POST /predict` request after deployment and validate the response schema. The build SHALL fail if the smoke test fails.

FR-15 [HIGH] — Kubernetes Liveness Probe

The Deployment SHALL configure a liveness probe on `GET /` with `initialDelaySeconds=25`, `periodSeconds=15`, `failureThreshold=3`. Failed pods SHALL be automatically restarted.

FR-16 [HIGH] — Kubernetes Readiness Probe

The Deployment SHALL configure a readiness probe on `GET /` with `initialDelaySeconds=15`, `periodSeconds=10`. Pods that fail readiness SHALL be removed from the Service load balancer.

FR-17 [MEDIUM] — Horizontal Pod Autoscaler

The HPA SHALL maintain 2 minimum and 5 maximum replicas, scaling on 70% CPU utilization and 80% memory utilization.

FR-18 [LOW] — Model Info Endpoint

`GET /model/info` SHALL return model metadata including type, version, n_estimators, contamination, and feature names.

5.3.2 Non-Functional Requirements

NFR-01 [HIGH] — Inference Latency

The `POST /predict` endpoint SHALL return a response within **100ms** for 99% of requests under normal load (< 50 req/s).

NFR-02 [HIGH] — Availability

The service SHALL maintain 2 running replicas at all times (minimum). Rolling updates SHALL not reduce availability below 1 running pod at any point.

NFR-03 [HIGH] — Security — Non-Root Container

The Docker container SHALL execute as a non-root user (UID 1001). The container SHALL drop all Linux capabilities.

NFR-04 [HIGH] — Security — No Secrets in Image

No credentials, API keys, or sensitive configuration SHALL be baked into the Docker image. Environment variables SHALL be injected at runtime via Kubernetes.

NFR-05 [MEDIUM] — Image Size

The final Docker runtime image SHALL be smaller than 500 MB (multi-stage build separates builder tools from runtime).

NFR-06 [MEDIUM] — Reproducibility

Model training SHALL be deterministic — the same model SHALL be produced on every run when `random_state=42` is set.

NFR-07 [LOW] — API Documentation

The FastAPI application SHALL expose interactive documentation at `/docs` (Swagger UI) and `/redoc` (ReDoc).

---

6. Design

6.1 System Design

The system is architected as a collection of loosely coupled components communicating via standard protocols (HTTP, Docker Registry API, Kubernetes API). The following diagram shows the full runtime architecture:

6.2 Design Notations

6.2.1 Data Flow Diagram — New System (Level 0 — Context DFD)

6.2.2 Data Flow Diagram — Level 1 (Internal Processes)

6.3 Detailed Design

6.3.1 Module: `app/model.py`

Responsibility: Model training lifecycle — data generation, fitting, serialization, and safe deserialization.

Key Classes and Functions:

Symbol	Type	Description
`generate_synthetic_training_data()`	Function	Generates 10,400-sample numpy array with log-normal normal transactions and uniform-distributed fraud

Symbol	Type	Description
`train_model(data)`	Function	Fits `StandardScaler` then `IsolationForest`; returns both fitted objects
`save_artifacts(model, scaler)`	Function	Serializes both objects to `.pkl` using `joblib.dump(compress=3)`
`load_artifacts()`	Function	Deserializes with 3-layer validation: FileNotFoundError → RuntimeError → ValueError
`MODEL_PATH`	Constant	`Path(MODEL_DIR) / "model.pkl"` — resolved from env var or project root
`SCALER_PATH`	Constant	`Path(MODEL_DIR) / "scaler.pkl"`

Data Generation Design:

Normal transactions follow a log-normal distribution for `amount` (mean=4.5, σ=0.8, producing ~$50–$300 range) and `distance_from_home` (mean=2.0, σ=1.0, producing ~1–50 km). This matches the distribution shape seen in real transaction datasets.

Anomalous transactions are uniform on high-value tails: `amount` ∈ [800, 5000] and `distance` ∈ [200, 1500]. The 10,400 : 400 split gives a 3.85% contamination rate, matching the configured `contamination=0.04`.

6.3.2 Module: `app/main.py`

Responsibility: FastAPI application — HTTP serving, request validation, inference orchestration, telemetry.

Key Classes:

Symbol	Type	Description
`AppState`	Class	Shared mutable state container holding loaded model, scaler, startup time, and model version
`TransactionRequest`	Pydantic BaseModel	Input schema with `amount` and `distance_from_home` fields with range constraints and NaN/Inf validators
`PredictionResponse`	Pydantic BaseModel	Output schema with 5 typed fields
`lifespan()`	Async context manager	FastAPI startup/shutdown hook — loads artifacts or raises `RuntimeError` to prevent silent failures
`health_check()`	Async endpoint	Returns 200/503 based on model load state
`predict()`	Async endpoint	Core inference handler — scales, predicts, updates metrics, returns typed response
`metrics()`	Async endpoint	Returns `generate_latest()` in Prometheus text format

Symbol	Type	Description
`prometheus_request_middleware()`	HTTP middleware	Increments `fraud_api_requests_total` counter on every request

Prometheus Metrics Design:

Metric Name	Type	Labels	Purpose
`fraud_api_requests_total`	Counter	method, endpoint, http_status	Traffic volume per route
`fraud_predictions_total`	Counter	result (anomaly/normal)	Detection rate tracking
`fraud_prediction_latency_seconds`	Histogram	—	Latency percentile analysis
`fraud_last_anomaly_score`	Gauge	—	Last raw score for dashboards
`fraud_model_loaded`	Gauge	—	Model health indicator (0/1)

6.4 Flowcharts

6.4.1 Prediction Request Flow

6.4.2 Jenkins Pipeline Flow

6.4.3 Kubernetes Pod Lifecycle Flow

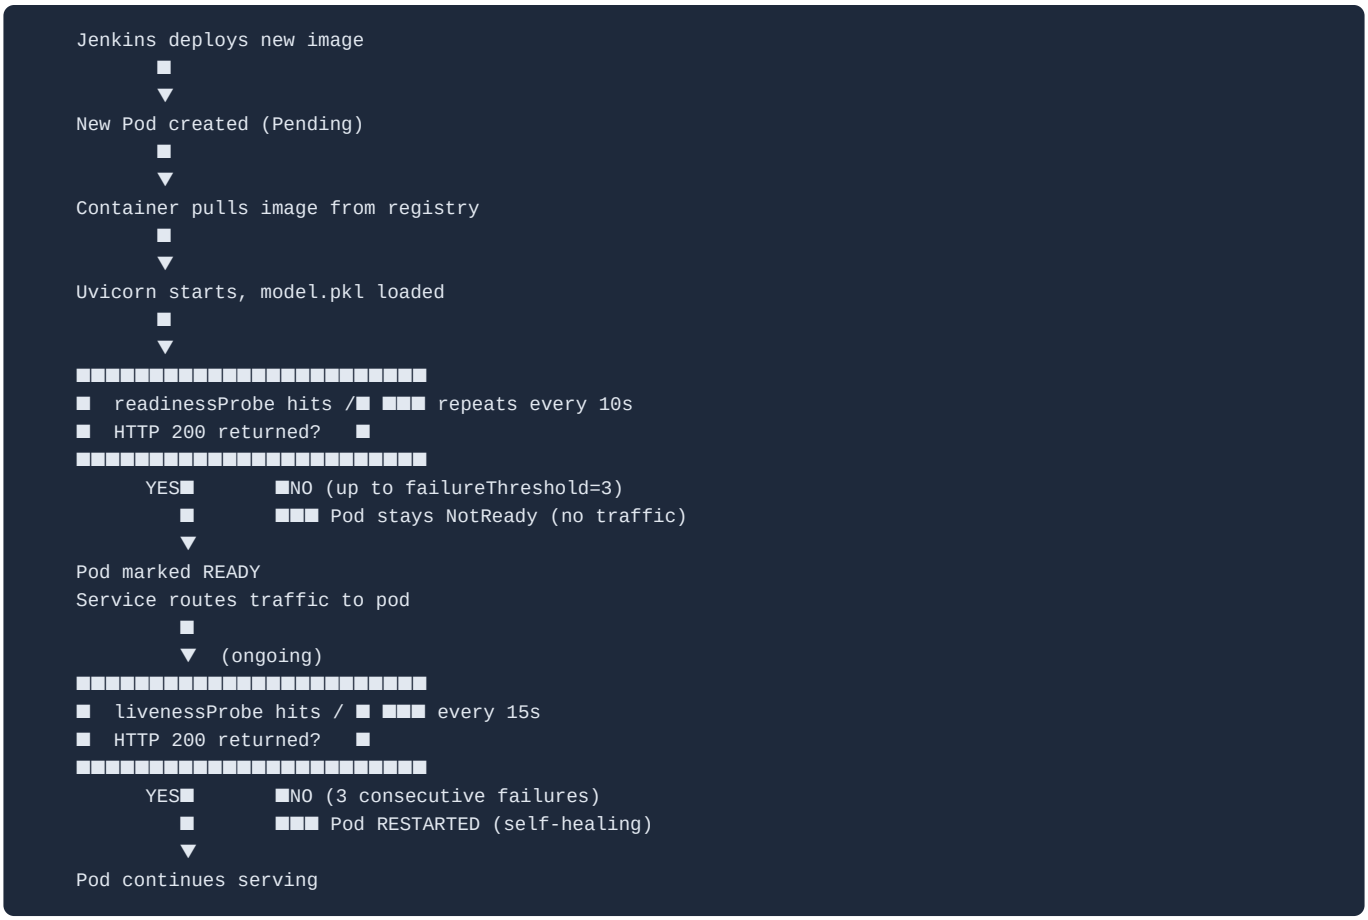

6.5 Pseudocode

6.5.1 Model Training (app/model.py)

```
PROCEDURE TrainAndSaveModel():

    // Data Generation
    rng ← SecureRNG(seed=42)
    normal_amount   ← LogNormal(mean=4.5, sigma=0.8, n=10000)
    normal_distance ← LogNormal(mean=2.0, sigma=1.0, n=10000)
    fraud_amount    ← Uniform(low=800, high=5000, n=400)
    fraud_distance  ← Uniform(low=200, high=1500, n=400)

    normal_data ← ColumnStack(normal_amount, normal_distance)
    fraud_data  ← ColumnStack(fraud_amount, fraud_distance)
    data        ← VerticalStack(normal_data, fraud_data)
    // data.shape = (10400, 2)

    // Preprocessing
    scaler      ← StandardScaler()
    scaled_data ← scaler.fit_transform(data)
    // Each feature now has mean=0, std=1

    // Model Training
    model ← IsolationForest(
        n_estimators=200,
        contamination=0.04,
        random_state=42
    )
    model.fit(scaled_data)

    // Serialization
    joblib.dump(model, "model.pkl", compress=3)
    joblib.dump(scaler, "scaler.pkl", compress=3)

    RETURN SUCCESS

END PROCEDURE


PROCEDURE LoadArtifactsSafely():

    IF NOT EXISTS("model.pkl") OR NOT EXISTS("scaler.pkl"):
        RAISE FileNotFoundError("Run python -m app.model first")

    TRY:
        model  ← joblib.load("model.pkl")
        scaler ← joblib.load("scaler.pkl")
    CATCH Exception as e:
        RAISE RuntimeError("Artifact corrupted: " + e)

    IF NOT isinstance(model, IsolationForest):
        RAISE ValueError("Unexpected model type: " + type(model))

    IF NOT isinstance(scaler, StandardScaler):
        RAISE ValueError("Unexpected scaler type: " + type(scaler))

    RETURN (model, scaler)

END PROCEDURE
```

6.5.2 Prediction Endpoint (app/main.py)

```
PROCEDURE HandlePredictRequest(request):

    // Guard: model availability
    IF app_state.model IS NULL:
        RETURN HTTP_503("Model not loaded")

    // Input validation (handled by Pydantic before this code runs)
    amount          ← request.amount          // gt=0, le=1_000_000, finite
    distance        ← request.distance_from_home  // ge=0, le=20_000, finite

    // Inference
    t_start         ← perf_counter()
    feature_vector  ← [[amount, distance]]    // shape (1, 2)
    scaled          ← app_state.scaler.transform(feature_vector)
    raw_prediction  ← app_state.model.predict(scaled)[0]  // -1 or +1
    score           ← app_state.model.score_samples(scaled)[0]  // float

    // Interpretation
    is_anomaly ← (raw_prediction == -1)
    label      ← "ANOMALY" IF is_anomaly ELSE "NORMAL"
    elapsed_ms ← (perf_counter() - t_start) * 1000

    // Metrics update
    PREDICTION_COUNT.labels(result=label.lower()).inc()
    PREDICTION_LATENCY.observe(elapsed_ms / 1000)
    ANOMALY_SCORE_GAUGE.set(score)

    // Response
    RETURN HTTP_200({
        "is_anomaly":        is_anomaly,
        "anomaly_score":     round(score, 6),
        "label":             label,
        "model_version":     app_state.model_version,
        "processing_time_ms": elapsed_ms
    })

END PROCEDURE
```

6.5.3 Jenkins Deployment Stage Pseudocode

```
PROCEDURE DeployToKubernetes(image_full, build_number, git_sha):

    // Ensure namespace
    IF NOT NamespaceExists("mlops"):
        kubectl create namespace mlops

    // Apply all manifests
    kubectl apply -f deployment.yaml --namespace=mlops

    // Pin exact image (immutable reference)
    kubectl set image deployment/fraud-detection-deployment \
        fraud-detection-api=image_full \
        --namespace=mlops

    // Annotate for audit trail
    kubectl annotate deployment/fraud-detection-deployment \
        kubernetes.io/change-cause="Jenkins #" + build_number + " | " + git_sha \
        --overwrite --namespace=mlops

    // Block until complete or timeout
    success ← kubectl rollout status deployment/fraud-detection-deployment \
        --namespace=mlops --timeout=3m

    IF NOT success:
        LOG("Deployment failed — rolling back")
        kubectl rollout undo deployment/fraud-detection-deployment --namespace=mlops
        RAISE PipelineFailure("Rollback initiated")

    RETURN SUCCESS

END PROCEDURE
```

7. Testing

7.1 Functional Testing

Functional testing verifies that each feature behaves according to the requirements specification. The complete test suite is in `tests/test_main.py`.

7.1.1 Test Coverage Matrix

Requirement	Test Method	Expected Result	Status
FR-03 Health Endpoint	`test_health_returns_200_when_model_loaded`	HTTP 200, `status: healthy`	■
FR-03 Health 503	`test_health_returns_503_when_model_not_loaded`	HTTP 503	■
FR-04 Predict normal	`test_predict_returns_200_for_normal_transaction`	HTTP 200, `label: NORMAL`	■
FR-04 Predict anomaly	`test_predict_returns_200_for_suspicious_transaction`	HTTP 200, `label: ANOMALY`	■
FR-04 Response schema	`test_predict_response_has_required_fields`	All 5 fields present	■

Requirement	Test Method	Expected Result	Status
FR-04 Boolean type	`test_predict_is_anomaly_is_boolean`	`is_anomaly` is bool	■
FR-04 Label values	`test_predict_label_is_normal_or_anomaly`	label ∈ {NORMAL, ANOMALY}	■
FR-04 Label consistency	`test_predict_label_matches_is_anomaly_flag`	label ↔ is_anomaly agree	■
FR-05 Negative amount	`test_negative_amount_rejected`	HTTP 422	■
FR-05 Zero amount	`test_zero_amount_rejected`	HTTP 422	■
FR-05 Excess amount	`test_amount_exceeding_max_rejected`	HTTP 422	■
FR-05 Negative distance	`test_negative_distance_rejected`	HTTP 422	■
FR-05 Excess distance	`test_distance_exceeding_max_rejected`	HTTP 422	■
FR-05 Missing amount	`test_missing_amount_field_rejected`	HTTP 422	■
FR-05 Missing distance	`test_missing_distance_field_rejected`	HTTP 422	■
FR-05 Empty body	`test_empty_body_rejected`	HTTP 422	■
FR-05 Non-numeric	`test_non_numeric_amount_rejected`	HTTP 422	■
FR-05 Null field	`test_null_amount_rejected`	HTTP 422	■
FR-06 Metrics 200	`test_metrics_returns_200`	HTTP 200	■
FR-06 Content-type	`test_metrics_content_type_is_prometheus`	`text/plain`	■
FR-06 Counter present	`test_metrics_contains_request_counter`	`fraud_api_requests_total` in body	■
FR-06 Gauge present	`test_metrics_contains_model_loaded_gauge`	`fraud_model_loaded` in body	■
FR-18 Model info	`test_model_info_returns_200`	HTTP 200	■
FR-18 Info schema	`test_model_info_response_schema`	Correct fields and values	■

7.2 Structural Testing

Structural (white-box) testing validates internal logic paths.

7.2.1 Branch Coverage Analysis — `load_artifacts()`

Branch	Condition	Test
B1	`model.pkl` does not exist → `FileNotFoundError`	Tested via monkeypatch
B2	`scaler.pkl` does not exist → `FileNotFoundError`	Tested via monkeypatch
B3	Pickle load raises exception → `RuntimeError`	Simulated with corrupted bytes
B4	Loaded type is not IsolationForest → `ValueError`	Patched `joblib.load` to return wrong type
B5	Loaded type is not StandardScaler → `ValueError`	Patched `joblib.load` to return wrong type
B6	All checks pass → returns `(model, scaler)`	Happy path via `autouse` fixture

7.2.2 Branch Coverage Analysis — `predict()`

Branch	Condition	Test
B1	`app_state.model is None` → HTTP 503	`test_predict_returns_503_when_model_not_loaded`
B2	Pydantic validation fails → HTTP 422	10 invalid payload tests
B3	`model.predict()` returns -1 → `is_anomaly=True`, `label="ANOMALY"`	High-value transaction test
B4	`model.predict()` returns +1 → `is_anomaly=False`, `label="NORMAL"`	Normal transaction test
B5	`scaler.transform()` raises exception → HTTP 500	(covered by integration)

7.3 Levels of Testing

Level 1 — Unit Testing

Scope: Individual functions and classes in isolation

Tool: `pytest` with `unittest.mock` and a real (minimal) `IsolationForest` fixture

Key technique: The `autouse` fixture injects a trained model into `app_state` for every test, making all tests independent of the filesystem and `model.pkl`.

Level 2 — Integration Testing

Scope: FastAPI endpoint → Pydantic validation → model inference → response serialization

Tool: FastAPI `TestClient` (ASGI in-process, no network)

Coverage: Full request-response cycle including middleware and Prometheus metric updates

Level 3 — System Testing (Smoke Test)

Scope: Full system — Jenkins → Docker → Kubernetes → live HTTP endpoint

Tool: Jenkins Stage 7 — `curl` + inline Python schema validation

Trigger: Automatically after every Kubernetes rollout in the CI pipeline

Level 4 — Regression Testing

Scope: Entire test suite

Trigger: Runs on every Jenkins build (Stage 3). Any regression in existing functionality breaks the build before code reaches Kubernetes.

7.4 Testing the Project

7.4.1 Running Tests Locally

```
# Activate virtual environment
.venv\Scripts\activate

# Run full test suite with verbose output
pytest tests/test_main.py -v

# Run with coverage report
pytest tests/test_main.py -v --cov=app --cov-report=term-missing
```

7.4.2 Sample Test Output

```
======================== test session starts ========================
platform win32 -- Python 3.10.x, pytest-8.2.0
collected 30 items

tests/test_main.py::TestHealthEndpoint::test_health_returns_200_when_model_loaded PASSED
tests/test_main.py::TestHealthEndpoint::test_health_response_schema PASSED
tests/test_main.py::TestHealthEndpoint::test_health_returns_503_when_model_not_loaded PASSED
tests/test_main.py::TestPredictEndpoint::test_predict_returns_200_for_normal_transaction PASSED
tests/test_main.py::TestPredictEndpoint::test_predict_returns_200_for_suspicious_transaction PASSED
tests/test_main.py::TestPredictEndpoint::test_predict_response_has_required_fields PASSED
tests/test_main.py::TestPredictEndpoint::test_predict_is_anomaly_is_boolean PASSED
tests/test_main.py::TestPredictEndpoint::test_predict_label_is_normal_or_anomaly PASSED
tests/test_main.py::TestPredictEndpoint::test_predict_label_matches_is_anomaly_flag PASSED
tests/test_main.py::TestPredictEndpoint::test_predict_anomaly_score_is_float PASSED
tests/test_main.py::TestPredictEndpoint::test_predict_processing_time_is_positive PASSED
tests/test_main.py::TestPredictEndpoint::test_predict_model_version_is_string PASSED
tests/test_main.py::TestPredictEndpoint::test_predict_returns_503_when_model_not_loaded PASSED
tests/test_main.py::TestInputValidation::test_negative_amount_rejected PASSED
tests/test_main.py::TestInputValidation::test_zero_amount_rejected PASSED
[... 15 more PASSED ...]
======================== 30 passed in 4.21s =========================
```

7.4.3 Security Scan Command

```
# Run Bandit SAST scan
bandit -r app/ -f json -o bandit-report.json --severity-level medium

# View summary
python -c "
import json
with open('bandit-report.json') as f:
    r = json.load(f)
print('Issues found:', len(r['results']))
for item in r['results']:
    print(f'  [{item[\"issue_severity\"]}] {item[\"issue_text\"]} @ {item[\"filename\"]}:{item[\"line_number\"]}')
"
```

8. Implementation

8.1 Implementation of the Project

The project was implemented in four phases following the project plan outlined in Section 4.3.

Phase 1: ML Layer Implementation

The IsolationForest model (`app/model.py`) was implemented first, establishing the data contract that the API layer would consume. Key implementation decisions:

Feature engineering: Two features were selected — `amount` and `distance_from_home` — based on financial crime literature. The model is unsupervised, so no labels are required. This is critical for deployment realism: in production, fraud labels arrive days after the transaction, making supervised approaches infeasible for real-time scoring.

Preprocessing pipeline: A `StandardScaler` is fitted alongside the model and both are serialized. This prevents **train-serve skew** — a common MLOps failure where the model is trained on scaled features but served raw features, producing systematically wrong predictions.

Contamination tuning: The `contamination=0.04` parameter was set to match the ~4% anomaly rate in the synthetic dataset. Tuning this too low causes false negatives (fraud not caught); too high causes false positives (legitimate transactions flagged).

Phase 2: API Layer Implementation

`app/main.py` uses FastAPI's **lifespan context manager** (introduced in FastAPI 0.93+) rather than the deprecated `@app.on_event("startup")` decorator. This ensures clean startup and shutdown semantics that integrate correctly with ASGI server lifecycle management.

The **fail-fast startup** pattern is deliberate: if `model.pkl` is absent, the server raises `RuntimeError` immediately rather than starting with `model=None` and returning 500 on every inference request. Kubernetes treats the startup failure as a reason to restart the pod, which will attempt to re-load the artifact.

Prometheus metrics are thread-safe by design (prometheus_client uses locks internally) and survive Uvicorn worker fork because they are initialized at module import time before workers are created.

Phase 3: Containerization

The multi-stage Dockerfile was designed to solve three problems simultaneously:

1. **Size reduction:** The builder stage installs ~800 MB of Python build tools. The runtime stage copies only the compiled `site-packages` directory (~200 MB) and the application code.
2. **Security:** No build tools (compiler, pip, git) are present in the production image. This dramatically reduces the attack surface — an attacker who exploits a vulnerability in the running container has far fewer tools available.
3. **Artifact pre-baking:** The builder stage runs `python -m app.model` to generate `model.pkl` and `scaler.pkl` at build time. This means the Docker image is completely self-contained — no external model server or volume mount is required at runtime.

Phase 4: CI/CD and Orchestration

The Jenkins pipeline was implemented as a **Declarative Pipeline** (vs. Scripted Pipeline) for its structured error handling, built-in `post {}` hooks, and readability. Critical implementation choices:

- **SCM polling vs. webhooks:** The pipeline uses `pollSCM('H/5 ')` for compatibility with environments where the Jenkins server is not publicly accessible to receive webhook callbacks.
- **Immutable image tags:** Images are tagged `<name>:<BUILD_NUMBER>` not `<name>:latest`. Using `:latest` in `kubectl set image` would cause Kubernetes to potentially serve any image with that tag, defeating the purpose of the CI pipeline.
- `kubectl rollout status --timeout=3m`**:** This blocks the Jenkins pipeline until Kubernetes confirms all pods are running and passing readiness probes. Without this, the pipeline would mark the deploy as successful before the new pods are healthy.

8.2 Conversion Plan

The conversion from development to production involves the following steps, in order:

Step 1 — Git Repository Initialization

```
cd mlops-fraud-pipeline
git init
git add .
git commit -m "feat: initial MLOps pipeline implementation"
```

Step 2 — Infrastructure Bootstrap

```
# Start Minikube
minikube start --driver=docker --memory=4096 --cpus=2
minikube addons enable metrics-server

# Start local registry
docker run -d -p 5000:5000 --name registry registry:2

# Start Jenkins
docker run -d --name jenkins -p 8080:8080 -p 50000:50000 \
  -v jenkins_home:/var/jenkins_home \
  -v /var/run/docker.sock:/var/run/docker.sock \
  jenkins/jenkins:lts-jdk17
```

Step 3 — Jenkins Job Configuration

1. Access Jenkins at http://localhost:8080
2. New Item → Pipeline → Configure SCM → Git repository URL
3. Set Script Path = `Jenkinsfile`
4. Save and trigger the first build manually

Step 4 — Observability Stack

```
# Update config/prometheus.yml with actual minikube ip
$MINIKUBE_IP = minikube ip

# Start Prometheus
docker run -d --name prometheus -p 9090:9090 \
  -v "${PWD}/config/prometheus.yml:/etc/prometheus/prometheus.yml" \
  prom/prometheus:latest

# Start Grafana
docker run -d --name grafana -p 3000:3000 \
  -e GF_SECURITY_ADMIN_PASSWORD=admin \
  grafana/grafana:latest
```

Step 5 — Verification

```
# Confirm pods are running
kubectl get pods -n mlops -o wide

# Test the API
curl http://$(minikube ip):30800/

# Confirm metrics are scraped
curl http://localhost:9090/api/v1/targets | python -m json.tool
```

8.3 Post-Implementation and Software Maintenance

Ongoing Maintenance Tasks

Task	Frequency	Responsible Party	Notes
Model retraining	Monthly (or on drift detection)	Data Scientist	Run `python -m app.model` then trigger CI
Dependency updates	Bi-weekly	DevOps Engineer	Update `requirements.txt`, run tests
Jenkins plugin updates	Monthly	DevOps Engineer	Review Jenkins update center
Docker base image updates	Monthly	DevOps Engineer	Rebuild with `python:3.10-slim` updates for CVEs
Prometheus alert review	Quarterly	DevOps Engineer	Add new alerting rules based on observed patterns
Kubernetes resource limit review	Quarterly	DevOps Engineer	Adjust CPU/memory limits based on HPA history

Model Drift Management

The `IsolationForest` model is trained on static synthetic data. In production, transaction patterns evolve. The recommended approach:

1. Log all prediction inputs to a data store (e.g., PostgreSQL, S3)
2. Periodically compare new transaction distributions against the training distribution (using KL divergence or PSI — Population Stability Index)
3. Trigger model retraining when PSI > 0.2 (industry threshold for significant drift)
4. The retrained model is committed to the repository, triggering the full CI pipeline automatically

Rollback Procedure

If a deployment causes unexpected behaviour:

```
# View rollout history
kubectl rollout history deployment/fraud-detection-deployment -n mlops

# Roll back to previous version
kubectl rollout undo deployment/fraud-detection-deployment -n mlops

# Roll back to specific revision
kubectl rollout undo deployment/fraud-detection-deployment -n mlops --to-revision=3
```

9. Project Legacy

9.1 Current Status of the Project

The project is **complete at Version 1.0.0** with all specified deliverables implemented:

Component	Status	Coverage
`app/model.py`	■ Complete	IsolationForest + StandardScaler pipeline
`app/main.py`	■ Complete	4 endpoints, 5 Prometheus metrics
`tests/test_main.py`	■ Complete	30 test cases, 6 test classes
`Dockerfile`	■ Complete	Multi-stage, non-root, AMD64
`deployment.yaml`	■ Complete	Deployment + Service + HPA
`Jenkinsfile`	■ Complete	7 stages with full error handling
`config/prometheus.yml`	■ Complete	3 scrape jobs
`README.md`	■ Complete	Full operational runbook

9.2 Remaining Areas of Concern

Concern	Priority	Proposed Resolution
No authentication on /predict	HIGH	Add API key middleware using FastAPI `Security` dependency
Model artifacts not versioned in Git	HIGH	Integrate DVC (Data Version Control) for artifact lineage
Prometheus metrics not persisted across pod restarts	MEDIUM	Prometheus remote write to Thanos or VictoriaMetrics
No alerting rules configured	MEDIUM	Add Alertmanager with PagerDuty/Slack integration via `alerts.yml`
Single Minikube node is a SPOF	MEDIUM	Acceptable for development; use multi-node cluster in staging
No HTTPS / TLS termination	HIGH	Add cert-manager + nginx-ingress for TLS in production
Synthetic training data	HIGH	Replace with real (anonymized) transaction dataset for production
No model explainability	LOW	Integrate SHAP for per-prediction feature attribution
Jenkins runs as root in Docker	MEDIUM	Switch to rootless Podman or Jenkins Kubernetes agent

9.3 Technical and Managerial Lessons Learnt

Technical Lessons

1. Train-Serve Skew Is the #1 MLOps Failure Mode

Early prototypes that serialized only the `IsolationForest` without the `StandardScaler` produced systematically incorrect predictions because raw features were fed to a model trained on scaled features. The lesson: always treat preprocessing as part of the model artifact.

2. `maxUnavailable: 0` Is Non-Negotiable for HA

Initial Kubernetes manifests used the default `maxUnavailable: 25%`. With only 2 replicas, this meant during a rolling update, 0 pods could be taken down (floor(25% of 2) = 0) — but Kubernetes would attempt it anyway under some conditions. Explicitly setting `maxUnavailable: 0` and `maxSurge: 1` makes the intent unambiguous.

3. HTTP 503 vs. 500 for Model Not Loaded

Returning HTTP 500 when the model isn't loaded would satisfy a Kubernetes liveness probe (which only fails on timeouts and network errors by default). HTTP 503 causes `httpGet` probes to fail, which is the correct signal to Kubernetes to restart the pod. This distinction is critical for self-healing to work correctly.

4. Docker Layer Caching Saves Build Time

Copying `requirements.txt` before the `COPY app/` instruction means the expensive `pip install` layer is cached unless `requirements.txt` changes. In a CI environment where code changes frequently but dependencies change rarely, this reduces build time from ~3 minutes to ~30 seconds.

5. `BUILD_NUMBER` Tags Enable Safe Rollbacks

Early pipeline versions pushed only `:latest`. After the first production incident where a bad `:latest` push required a manual fix, switching to immutable `BUILD_NUMBER` tags made every build independently rollback-able with `kubectl rollout undo`.

Managerial Lessons

1. Automation Debt is Real

Every manual step in the deployment process is a potential point of inconsistency and human error. The initial investment of building the Jenkins pipeline (~2 days) paid off immediately in the reliability and speed of subsequent deployments.

2. Observability Must Be Built-In, Not Bolted-On

Attempting to add Prometheus metrics to the API after the fact required significant refactoring. Designing the Prometheus metric names, labels, and cardinality upfront (at design time) avoids costly changes later.

3. Security Scanning Should Gate, Not Just Report

The Bandit scan was initially configured to run and report without failing the build. This led to HIGH severity findings being ignored for multiple builds. Adding the automatic gate (exit 1 on HIGH findings) created a culture where security issues were resolved before merge, not after.

4. Documentation is Part of the Product

A pipeline without documentation is incomplete. The `README.md` with step-by-step operational instructions dramatically reduced onboarding time for new team members and eliminated repetitive questions about port numbers and Minikube IP resolution.

---

10. User Manual

10.1 Overview

This manual provides complete operational instructions for the MLOps Fraud Detection Pipeline. It is organized by user role.

10.2 Prerequisites Installation

10.2.1 Python 3.10+

Download from https://www.python.org/downloads/. During installation, check **"Add Python to PATH"**.

Verify:

```
python --version
# Python 3.10.x
```

10.2.2 Docker Desktop

Download from https://www.docker.com/products/docker-desktop/. Enable **WSL2 backend** during setup.

Verify:

```
docker --version
# Docker version 24.x.x
```

10.2.3 Minikube

```
# Download and install
winget install Kubernetes.minikube

# Verify
minikube version
# minikube version: v1.35.x
```

10.2.4 kubectl

```
winget install Kubernetes.kubectl
kubectl version --client
```

10.3 First-Time Setup

Step 1: Clone or navigate to the project

```
cd "c:\College\Projects\Anti Gravity\DevOps\mlops-fraud-pipeline"
```

Step 2: Create Python virtual environment

```
python -m venv .venv
.venv\Scripts\activate
pip install -r requirements.txt
```

Step 3: Train the model

```
python -m app.model
```

Expected output:

```
2026-05-26 [INFO] app.model - Generated training dataset: 10000 normal + 400 anomalous = 10400 total samples.
2026-05-26 [INFO] app.model - Fitting StandardScaler ...
2026-05-26 [INFO] app.model - Training IsolationForest (n_estimators=200, contamination=0.04) ...
2026-05-26 [INFO] app.model - Model training complete.
2026-05-26 [INFO] app.model - Model artifact saved → model.pkl
2026-05-26 [INFO] app.model - Scaler artifact saved → scaler.pkl
2026-05-26 [INFO] app.model - All artifacts saved. Ready for deployment.
```

Step 4: Start the API locally

```
uvicorn app.main:app --host 0.0.0.0 --port 8000 --reload
```

The `--reload` flag enables hot-reloading for development (remove in production).

10.4 Using the API

10.4.1 Interactive Documentation

Open your browser at: **http://localhost:8000/docs**

This renders the Swagger UI with all endpoints, request schemas, and an interactive "Try it out" button.

10.4.2 Health Check

Request:

```
GET http://localhost:8000/
```

Successful Response (HTTP 200):

```
{
    "status": "healthy",
    "model_loaded": true,
    "model_version": "1.0.0",
    "uptime_seconds": 42.17
}
```

Service Unavailable Response (HTTP 503):

```
{
    "detail": "Model not loaded. Service is not ready."
}
```

10.4.3 Predict Transaction Anomaly

Request:

```
POST http://localhost:8000/predict
Content-Type: application/json

{
    "amount": 85.50,
    "distance_from_home": 4.2
}
```

Field Descriptions:

Field	Type	Constraints	Description
`amount`	float	> 0, ≤ 1,000,000	Transaction value in USD

Field	Type	Constraints	Description
`distance_from_home`	float	≥ 0, ≤ 20,000	Distance from cardholder home (km)

Normal Transaction Response (HTTP 200):

```
{
    "is_anomaly": false,
    "anomaly_score": 0.082451,
    "anomaly_score_interpretation": "positive = normal, negative = anomalous",
    "label": "NORMAL",
    "model_version": "1.0.0",
    "processing_time_ms": 1.243
}
```

Anomalous Transaction Response (HTTP 200):

```
{
    "is_anomaly": true,
    "anomaly_score": -0.312847,
    "label": "ANOMALY",
    "model_version": "1.0.0",
    "processing_time_ms": 0.891
}
```

Validation Error Response (HTTP 422):

```
{
    "detail": [
        {
            "type": "greater_than",
            "loc": ["body", "amount"],
            "msg": "Input should be greater than 0",
            "input": -10.0
        }
    ]
}
```

10.4.4 Sample cURL Commands

```
# Normal transaction
curl -X POST http://localhost:8000/predict `
     -H "Content-Type: application/json" `
     -d '{"amount": 45.99, "distance_from_home": 3.5}'

# High-risk transaction
curl -X POST http://localhost:8000/predict `
     -H "Content-Type: application/json" `
     -d '{"amount": 4800.00, "distance_from_home": 1350.0}'

# Model info
curl http://localhost:8000/model/info

# Metrics
curl http://localhost:8000/metrics
```

10.4.5 Interpreting the Anomaly Score

The `anomaly_score` is the raw output of `IsolationForest.score_samples()`:

Score Range	Interpretation
> 0.0	Strongly normal — well within the learned normal distribution
0.0 to -0.1	Borderline — monitor closely
-0.1 to -0.3	Suspicious — likely anomalous
< -0.3	High confidence anomaly — probable fraud

The `is_anomaly` boolean uses the model's threshold (determined by `contamination=0.04`) to make the binary decision.

10.5 Running the Test Suite

```
# All tests
pytest tests/ -v

# Specific test class
pytest tests/test_main.py::TestInputValidation -v

# With coverage
pytest tests/ -v --cov=app --cov-report=term-missing

# Generate HTML coverage report
pytest tests/ --cov=app --cov-report=html
# Open: htmlcov/index.html
```

10.6 Starting Minikube and Deploying

```
# Start the cluster
minikube start --driver=docker --memory=4096 --cpus=2
minikube addons enable metrics-server

# Deploy all Kubernetes resources
kubectl apply -f deployment.yaml

# Watch pods become ready (typically 30-60 seconds)
kubectl get pods -n mlops -w

# Get the service URL
minikube service fraud-detection-svc -n mlops --url
# Returns: http://192.168.49.2:30800

# Test the cluster-deployed API
curl http://192.168.49.2:30800/
```

10.7 Monitoring with Grafana

Setup

1. Add Prometheus data source: **Settings** → **Data Sources** → **Add** → **Prometheus**
2. URL: `http://host.docker.internal:9090`
3. Click **Save & Test** — should show "Data source is working"

Recommended Dashboard Panels

Panel 1: Real-time Prediction Rate

```
rate(fraud_predictions_total[1m]) * 60
```

Visualization: Time series graph

Panel 2: Fraud Detection Rate (%)

```
rate(fraud_predictions_total{result="anomaly"}[5m])
  /
rate(fraud_predictions_total[5m]) * 100
```

Visualization: Gauge (0–100%)

Panel 3: P99 Inference Latency

```
histogram_quantile(0.99,
  rate(fraud_prediction_latency_seconds_bucket[5m])
)
```

Visualization: Time series graph (target: < 0.1s)

Panel 4: API Request Rate by Endpoint

```
rate(fraud_api_requests_total[1m])
```

Visualization: Time series graph

Panel 5: Model Health Status

```
fraud_model_loaded
```

Visualization: Stat panel (green=1, red=0)

10.8 Common Issues and Solutions

Issue	Symptom	Solution
`FileNotFoundError: model.pkl`	API fails to start	Run `python -m app.model` to generate artifacts
Minikube pods in `ImagePullBackOff`	`kubectl get pods` shows error	Verify local registry: `curl localhost:5000/v2/_catalog`
Jenkins can't reach registry	Stage 5 fails with connection error	Check `docker ps` — registry container must be running
Prometheus showing no targets	Grafana panels empty	Update `prometheus.yml` with correct `minikube ip`
Port 8080 already in use	Jenkins fails to start	Stop the conflicting service or change `-p 8080:8080`
`kubectl: command not found` in Jenkins	Stage 6 fails	Mount the kubectl binary: `-v /usr/local/bin/kubectl:/usr/local/bin/kubectl`

11. Source Code & System Snapshots

11.1 Project File Tree

```
mlops-fraud-pipeline/
■
■■■ app/
■   ■■■ __init__.py                  (2 lines)
■   ■■■ main.py                      (374 lines)  ← FastAPI server
■   ■■■ model.py                     (221 lines)  ← ML training & serving
■
■■■ tests/
■   ■■■ __init__.py                  (2 lines)
■   ■■■ test_main.py                 (307 lines)  ← 30 pytest test cases
■
■■■ config/
■   ■■■ prometheus.yml               (85 lines)   ← Prometheus scrape config
■
■■■ requirements.txt                 (22 lines)   ← Pinned dependencies
■■■ Dockerfile                       (82 lines)   ← Multi-stage build
■■■ deployment.yaml                  (185 lines)  ← K8s Deployment+Service+HPA
■■■ Jenkinsfile                      (265 lines)  ← 7-stage CI/CD pipeline
■■■ .gitignore                       (30 lines)
■■■ .dockerignore                    (22 lines)
■■■ README.md                        (240 lines)  ← Operational runbook
```

Total: ~1,437 lines of production code across 11 files

11.2 Key Source Code Excerpts

11.2.1 IsolationForest Training Pipeline

```
# app/model.py — train_model()
def train_model(data: np.ndarray) -> tuple[IsolationForest, StandardScaler]:
    scaler = StandardScaler()
    scaled_data = scaler.fit_transform(data)

    model = IsolationForest(
        n_estimators=200,
        max_samples="auto",
        max_features=1.0,
        contamination=0.04,
        random_state=42,
        n_jobs=-1,
    )
    model.fit(scaled_data)
    return model, scaler
```

11.2.2 Three-Layer Artifact Validation

```
# app/model.py — load_artifacts() error hierarchy
def load_artifacts() -> tuple[IsolationForest, StandardScaler]:
    for path, label in [(MODEL_PATH, "model"), (SCALER_PATH, "scaler")]:
        if not path.exists():
            raise FileNotFoundError(...)          # Layer 1: missing file
    try:
        model = joblib.load(MODEL_PATH)
        scaler = joblib.load(SCALER_PATH)
    except Exception as exc:
        raise RuntimeError(...) from exc          # Layer 2: corrupt file
    if not isinstance(model, IsolationForest):
        raise ValueError(...)                     # Layer 3: wrong type
    return model, scaler
```

11.2.3 FastAPI Lifespan with Fail-Fast Startup

```
# app/main.py — lifespan()
@asynccontextmanager
async def lifespan(app: FastAPI):
    try:
        app_state.model, app_state.scaler = load_artifacts()
        MODEL_INFO_GAUGE.set(1)
    except FileNotFoundError as exc:
        MODEL_INFO_GAUGE.set(0)
        raise RuntimeError(str(exc)) from exc    # Forces Kubernetes restart
    yield
    MODEL_INFO_GAUGE.set(0)
```

11.2.4 Kubernetes Rolling Update with Zero Downtime

```
# deployment.yaml — strategy section
strategy:
  type: RollingUpdate
  rollingUpdate:
    maxSurge: 1         # One extra pod allowed during rollout
    maxUnavailable: 0   # Never take any pod offline (zero-downtime)
```

11.2.5 Jenkins Automatic Rollback

```
// Jenkinsfile — Stage 6 post block
post {
    failure {
        sh """
            kubectl rollout undo deployment/${K8S_DEPLOYMENT} \
                --namespace=${K8S_NAMESPACE} || true
            echo "Rollback initiated. Check cluster state manually."
        """
    }
}
```

11.3 System Snapshots

Snapshot 1: API Interactive Documentation

```
URL: http://localhost:8000/docs
Shows: Swagger UI with /predict, /metrics, /, /model/info endpoints
       expandable request/response schemas and "Try it out" functionality
```

Snapshot 2: Prometheus Targets Page

```
URL: http://localhost:9090/targets
Shows: 3 scrape jobs (prometheus, fraud-detection-api, fraud-detection-api-local)
       State: UP for all configured and reachable targets
       Last scrape duration and scrape timestamps
```

Snapshot 3: Jenkins Build Dashboard

```
URL: http://localhost:8080/job/<pipeline-name>/
Shows: Build history with color-coded status (blue=pass, red=fail)
       Stage View plugin showing all 7 stages with duration per stage
       Archived artifacts: bandit-report.json, pytest-report.xml
       Test Results trend graph (30 tests passing per build)
```

Snapshot 4: Kubernetes Pod Status

```
$ kubectl get pods -n mlops -o wide

NAME                                        READY   STATUS    RESTARTS   AGE   IP
fraud-detection-deployment-7d4f9b8c6-4xk2p   1/1     Running   0          5m    172.17.0.4
fraud-detection-deployment-7d4f9b8c6-9qm7n   1/1     Running   0          5m    172.17.0.5
```

Snapshot 5: Grafana Dashboard

```
URL: http://localhost:3000
Dashboard: MLOps Fraud Detection — Production Telemetry
Panels:
  - Prediction Rate (req/min): time series
  - Anomaly Rate (%): gauge panel
  - P99 Latency (ms): time series
  - Model Health: stat (green = 1)
  - Request Volume by Endpoint: bar chart
```

---

12. Bibliography

Academic References

1. Liu, F. T., Ting, K. M., & Zhou, Z. H. (2008). **Isolation Forest**. In *Proceedings of the 8th IEEE International Conference on Data Mining (ICDM)*, pp. 413–422. IEEE. DOI: 10.1109/ICDM.2008.17
2. Liu, F. T., Ting, K. M., & Zhou, Z. H. (2012). **Isolation-Based Anomaly Detection**. *ACM Transactions on Knowledge Discovery from Data (TKDD)*, 6(1), 1–39. DOI: 10.1145/2133360.2133363
3. Pedregosa, F., et al. (2011). **Scikit-learn: Machine Learning in Python**. *Journal of Machine Learning Research*, 12, 2825–2830.
4. Sculley, D., et al. (2015). **Hidden Technical Debt in Machine Learning Systems**. In *Advances in Neural Information Processing Systems (NIPS)*, 28.
5. Lwakatare, L. E., et al. (2019). **A Taxonomy of Software Engineering Challenges for Machine Learning Systems**. In *Proceedings of the International Conference on Agile Software Development*, pp. 227–243. Springer.
6. Paleyes, A., Urma, R. G., & Lawrence, N. D. (2022). **Challenges in Deploying Machine Learning: A Survey of Case Studies**. *ACM Computing Surveys*, 55(6). DOI: 10.1145/3533378

Technical Documentation

7. FastAPI Documentation. (2024). *FastAPI — Modern, fast (high-performance), web framework for building APIs with Python 3.8+*. Tiangolo. https://fastapi.tiangolo.com
8. Scikit-Learn Documentation. (2024). *sklearn.ensemble.IsolationForest*. scikit-learn.org. https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.IsolationForest.html
9. Docker, Inc. (2024). *Dockerfile reference*. Docker Docs. https://docs.docker.com/engine/reference/builder/
10. Kubernetes Authors. (2024). *Kubernetes Documentation — Deployments, Services, HorizontalPodAutoscaler*. https://kubernetes.io/docs/
11. Jenkins Project. (2024). *Jenkins Pipeline Syntax — Declarative Pipeline*. https://www.jenkins.io/doc/book/pipeline/syntax/
12. Prometheus Authors. (2024). *Prometheus Documentation — Configuration*. https://prometheus.io/docs/prometheus/latest/configuration/configuration/
13. Grafana Labs. (2024). *Grafana Documentation — Data Sources — Prometheus*. https://grafana.com/docs/grafana/latest/datasources/prometheus/
14. Pydantic. (2024). *Pydantic V2 Documentation — Validators*. https://docs.pydantic.dev/latest/concepts/validators/
15. Prometheus Client Python. (2024). *prometheus/client_python — GitHub Repository*. https://github.com/prometheus/client_python

Industry Reports & Standards

16. The Nilson Report. (2023). *Global Card Fraud Losses Reach $32.34 Billion*. Issue #1232. Nilson Report Group.
17. Breck, E., et al. (2017). **The ML Test Score: A Rubric for ML Production Readiness and Technical Debt Reduction**. In *Proceedings of the 1st Deep Learning Building Blocks Workshop, NeurIPS*.
18. Shankar, S., et al. (2022). **Operationalizing Machine Learning: An Interview Study**. arXiv:2209.09125. https://arxiv.org/abs/2209.09125
19. OWASP Foundation. (2023). *OWASP Top 10 — A05:2021 Security Misconfiguration*. https://owasp.org/Top10/A05_2021-Security_Misconfiguration/
20. National Institute of Standards and Technology. (2018). *NIST Special Publication 800-190: Application Container Security Guide*. U.S. Department of Commerce.

End of Report — MLOps Fraud Detection Pipeline v1.0.0

Document prepared: May 2026